

h1ybyyvxx

March 30, 2023

1. How does similarity is calculated if data is categorical in nature.

Categorical data (also known as nominal data) has been studied for a long time in various contexts. However computing similarity between categorical data instances is not straightforward owing to the fact that there is no explicit notion of ordering between categorical values. To overcome this problem, several data-driven similarity measures have been proposed for categorical data. The behavior of such measures directly depends on the data.

Let us try to get to know some of the measures. To define these measures lets begin with a few notations.

Definitions :

Consider a categorical data set D containing N data points (rows), defined over a set of d categorical attributes let A_k represent the k th attribute where $K=1,2,3\dots d$. Let A_k take n_k unique values in the data set. Let $f_k(x)$ be the frequency of the attribute A_k taking the value x which is one of the n_k values. let $p_k(x)$ be the sample probability of the attribute A_k taking the value x . There are many similarity measures in the literature here we are only going to define a few of them and implement them in python.

The relationship between similarity measure and distance measure :

$\text{sim} = 1/\text{len}$ where sim = similarity and len = distance. Similarity measures assigns value between two data instances X and Y belonging to the data set D as follows:

$S(X,Y) = \sum_k w_k S_k(X,Y)$ Here w_k is weight assigned to each attribute and $S_k(X,Y)$ is a similarity score function for k th attribute in the given two instances of the data set i.e, X and Y . With different definitions for w_k and $S_k(X,Y)$ we get different similarity measures having different properties and use cases. Here all the measures are defined using w_k and $S_k(X,Y)$

Overlap Definition : $S_k(X,Y) = 1$ if $X_k = Y_k$ and equal to 0 otherwise and weight w_k is $1/d$ here $k = 1,2,\dots,d$

Eskin Definition: $S_k(X,Y) = 1$ if $X_k = Y_k$ and equal to $n_k^2/(n_k^2+2)$ otherwise and the weight w_k is $1/d$ here $k = 1,2,\dots,d$

IOF-Inverse Occurrence Frequency Definition: $S_k(X,Y) = 1$ if $X_k = Y_k$ and equal to

$1/(1+\log(f_k(X_k))*\log(f_k(Y_k)))$ otherwise and the weight w_k is $1/d$ here $k = 1,2,\dots,d$

OF Definition: $S_k(X,Y) = 1$ if $X_k = Y_k$ and equal to

$1/(1+\log(N/f_k(X_k))*\log(N/f_k(Y_k)))$ otherwise and the weight w_k is $1/d$ here $k = 1,2,\dots,d$

Lin Definition: $S(X, Y) = 2\log(p(X))$ if $X = Y$ and equal to $2\log(p(X) + p(Y))$ otherwise and the weight w is $1/(\log(p(X)) + \log(p(Y)))$ here $k = 1, 2, \dots, d$

2. Implement K-means for “Car dataset” and come up with the business insights

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: data = pd.read_csv('Cars_mileage.csv')
```

```
[3]: data.head()
```

```
[3]:
```

	HP	MPG	VOL	SP	WT
0	49	53.700681	89	104.185353	28.762059
1	55	50.013401	92	105.461264	30.466833
2	55	50.013401	92	105.461264	30.193597
3	70	45.696322	92	113.461264	30.632114
4	53	50.504232	92	104.461264	29.889149

```
[4]: data.describe()
```


```
[4]:
```

	HP	MPG	VOL	SP	WT
count	81.000000	81.000000	81.000000	81.000000	81.000000
mean	117.469136	34.422076	98.765432	121.540272	32.412577
std	57.113502	9.131445	22.301497	14.181432	7.492813
min	49.000000	12.101263	50.000000	99.564907	15.712859
25%	84.000000	27.856252	89.000000	113.829145	29.591768
50%	100.000000	35.152727	101.000000	118.208698	32.734518
75%	140.000000	39.531633	113.000000	126.404312	37.392524
max	322.000000	53.700681	160.000000	169.598513	52.997752

```
[5]: sns.boxplot(data['HP'])
```

```
C:\Users\87548\anaconda3\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
warnings.warn(
```

```
[5]: <AxesSubplot:xlabel='HP'>
```



```
[6]: data.columns
```

```
[6]: Index(['HP', 'MPG', 'VOL', 'SP', 'WT'], dtype='object')
```

```
[7]: # Removing Outliers
from scipy import stats

z = np.abs(stats.zscore(data[['HP', 'MPG', 'VOL', 'SP', 'WT']]))

print(z)

print('*****')

print(np.where(z > 3))
```

```
[[1.20629511  2.12438703  0.44061061  1.23140253  0.49023816]
 [1.1005866   1.7180708   0.30525236  1.14087145  0.26129934]
 [1.1005866   1.7180708   0.30525236  1.14087145  0.29799301]
 [0.83631531  1.24235454  0.30525236  0.57323872  0.2391033 ]
 [1.13582277  1.77215741  0.30525236  1.21182554  0.33887816]
 [0.83631531  1.24235454  0.44061061  0.59281571  0.37881419]
 [1.1005866   1.7180708   0.30525236  1.14087145  0.28256505]
 [0.97726     1.354778    2.20026792  1.34399532  2.22453536]
 [0.97726     1.354778    2.20026792  1.34399532  2.15581433]
 [0.66013445  0.86799855  0.21501352  0.4182792   0.20042158]]
```

[0.78346105 1.12736839 0.44061061 0.73472389 0.40949027]
 [0.44871742 0.54347892 2.20026792 0.27968395 2.23718888]
 [0.44871742 0.54347892 0.01058357 0.04007367 0.05385348]
 [0.78346105 1.12736839 0.44061061 0.73472389 0.40746314]
 [0.90678765 1.24660479 0.44061061 0.94758617 0.41164734]
 [0.78346105 1.12736839 0.44061061 0.73472389 0.37710085]
 [0.69537062 0.92208516 0.35037178 0.50881029 0.38633253]
 [0.44871742 0.54347892 2.20026792 0.27968395 2.17802125]
 [0.69537062 0.92208516 0.35037178 0.50881029 0.33347355]
 [0.48395359 0.49364256 0.19106124 0.21764002 0.14831256]
 [0.44871742 0.43955596 0.01058357 0.1727886 0.01180799]
 [0.76584297 0.9263354 0.37153891 0.75917009 0.33514564]
 [0.39586316 0.42849277 0.1008224 0.08878318 0.03535268]
 [0.64251637 0.66696556 0.12477469 0.54713605 0.07727938]
 [0.39586316 0.42849277 0.44061061 0.16709116 0.48759676]
 [0.44871742 0.43955596 2.20026792 0.49254623 2.19829229]
 [0.44871742 0.43955596 0.82273309 0.05532661 0.75878723]
 [0.44871742 0.43955596 0.01058357 0.1727886 0.05673768]
 [1.15344086 0.99696475 0.23618066 1.55924209 0.27807071]
 [0.25491848 0.10822339 0.37153891 0.02132491 0.42124726]
 [0.58966211 0.55197941 0.68737484 0.57158225 0.62174504]
 [0.58966211 0.55197941 0.1008224 0.65641591 0.11035971]
 [0.27253656 0.20533341 0.07965527 0.11488584 0.13866213]
 [0.27253656 0.20533341 0.64225542 0.01047519 0.69304639]
 [0.64251637 0.5630426 0.1008224 0.72737 0.03881962]
 [0.48395359 0.3897196 0.03453585 0.46313063 0.06732815]
 [0.48395359 0.3897196 0.48573003 0.52838729 0.49132041]
 [0.27253656 0.038723 0.57596887 0.32857635 0.60870364]
 [0.27253656 0.038723 0.57596887 0.32857635 0.50790524]
 [0.22076984 0.37553513 0.30525236 0.49107265 0.30848967]
 [0.39586316 0.0805134 0.64225542 0.36524564 0.66877093]
 [0.39586316 0.0805134 0.32641949 0.4109253 0.35112169]
 [0.27253656 0.038723 0.30525236 0.28942235 0.253162]
 [0.39586316 0.0805134 0.48573003 0.52838729 0.54638368]
 [0.43109934 0.1346 0.14594182 0.50798206 0.08944633]
 [0.30777273 0.01536361 0.01058357 0.31469678 0.02811271]
 [0.30777273 0.01536361 0.55201658 0.23638879 0.54872568]
 [0.34300891 0.06945021 0.19106124 0.35954821 0.20110061]
 [0.22076984 0.37553513 0.57596887 0.45191866 0.58309656]
 [0.04350145 0.52806514 0.1008224 0.23069136 0.14045608]
 [0.04350145 0.52806514 0.1008224 0.23069136 0.10761914]
 [0.04350145 0.52806514 0.1008224 0.23069136 0.13753382]
 [0.04350145 0.52806514 1.13856902 0.08060104 1.07240653]
 [1.10167413 1.09474551 0.64225542 1.55051482 0.69941256]
 [0.74931242 0.83414636 0.64225542 0.98288209 0.65022074]
 [0.22076984 0.72351451 1.13856902 0.3451235 1.09804367]
 [0.37824508 0.36457601 0.30525236 0.78610099 0.30417763]
 [0.04350145 0.52806514 0.1008224 0.23069136 0.04323438]

```
[0.30777273 0.47274922 0.21501352 0.63114147 0.24136353]
[0.30777273 0.61288266 0.73249426 0.4231484 0.70507688]
[0.48504113 0.77883043 0.55201658 0.6150603 0.60103751]
[0.04458898 1.08132775 0.77761368 0.27471455 0.73160477]
[0.3969507 1.20179346 1.45440495 0.31984907 1.47432127]
[0.3969507 1.20179346 1.0934496 0.26764375 1.11600704]
[0.57313156 1.1919596 1.00321077 0.46745469 1.0403533 ]
[0.83740285 0.62016359 2.20026792 0.35890287 2.24265139]
[0.83740285 1.24727552 0.68737484 0.77654549 0.74765875]
[0.83740285 1.24727552 1.27392727 0.86137914 1.23030947]
[0.83740285 1.24727552 1.0934496 0.83527648 1.08232792]
[2.24684972 1.44887149 0.597136 2.60830053 0.63509149]
[2.86348272 1.62465312 2.20026792 3.05515833 2.22785207]
[0.78454859 1.23621234 1.63488262 0.84263038 1.55787152]
[0.78454859 1.23621234 1.49952437 0.82305338 1.4692362 ]
[0.3969507 1.68990629 2.76286807 0.22527702 2.76444017]
[0.3969507 1.68990629 1.36416611 0.02298138 1.3706083 ]
[1.01358371 1.72555448 1.36416611 0.80347638 1.39203071]
[3.60344233 0.27305243 2.20026792 3.40992879 2.18623649]
[2.12352311 1.67761447 0.73249426 2.06024479 0.74002517]
[2.56397526 0.04651022 2.20026792 2.13275515 2.23502801]
[3.12775401 1.60754809 0.91297193 3.292567 0.94146241]
[2.08828694 2.45962006 0.37153891 1.29849856 0.34057153]]
```

```
*****
```

```
(array([70, 76, 76, 79, 79], dtype=int64), array([3, 0, 3, 0, 3], dtype=int64))
```

```
[8]: data_outlier_removed = data[(z<3).all(axis=1)]
```

```
[9]: data_outlier_removed.head(10)
```

```
[9]:   HP      MPG  VOL      SP      WT
0  49  53.700681   89  104.185353  28.762059
1  55  50.013401   92  105.461264  30.466833
2  55  50.013401   92  105.461264  30.193597
3  70  45.696322   92  113.461264  30.632114
4  53  50.504232   92  104.461264  29.889149
5  70  45.696322   89  113.185353  29.591768
6  55  50.013401   92  105.461264  30.308480
7  62  46.716554   50  102.598513  15.847758
8  62  46.716554   50  102.598513  16.359484
9  80  42.299078   94  115.645204  30.920154
```

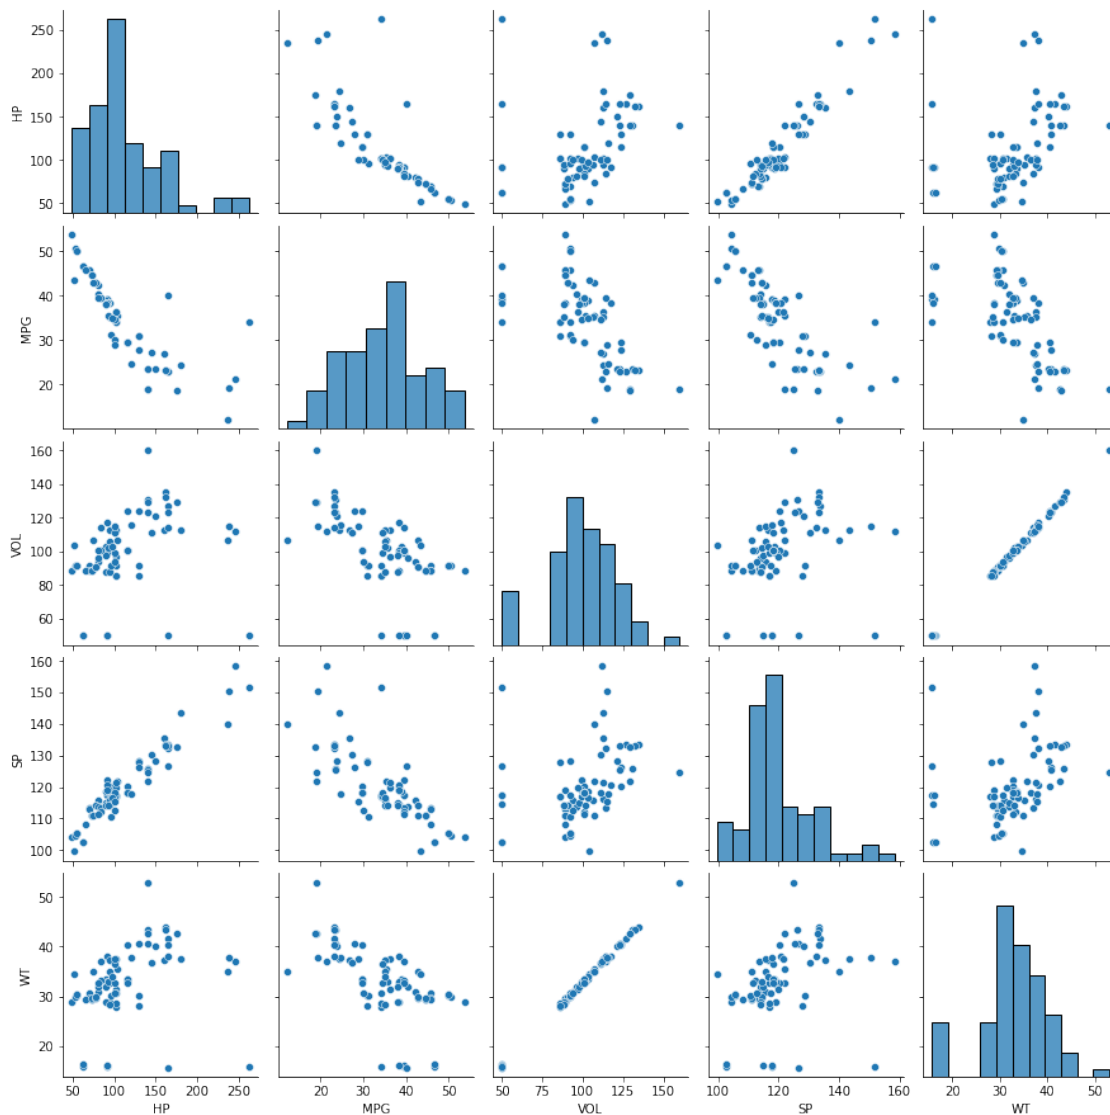
```
[10]: print('Shape of dataframe before outlier removal: ' + str(data.shape))
      print('Shape of dataframe after outlier removal: ' + str(data_outlier_removed.
      ↪shape))
```

```
Shape of dataframe before outlier removal: (81, 5)
```

```
Shape of dataframe after outlier removal: (78, 5)
```

```
[11]: sns.pairplot(data_outlier_removed)
```

```
[11]: <seaborn.axisgrid.PairGrid at 0x293fe6f5f40>
```



```
[13]: x = data_outlier_removed
```

```
[14]: from sklearn.cluster import KMeans
wcsc = []


for i in range(1,11):
    kmeans = KMeans(
        n_clusters=i,init='k-means++',max_iter=300,n_init=10,random_state=0)
    kmeans.fit(x)
```

```
wcss.append(kmeans.inertia_)

plt.plot(range(1,11),wcss)
plt.title('Elbow Method')
plt.show()
```

C:\Users\87548\anaconda3\lib\site-packages\sklearn\cluster_kmeans.py:881:
UserWarning: KMeans is known to have a memory leak on Windows with MKL, when
there are less chunks than available threads. You can avoid it by setting the
environment variable OMP_NUM_THREADS=1.

```
warnings.warn(
```



```
[15]: kmeans = KMeans(n_clusters=4,init='k-means++',max_iter=300,n_init=10,random_state=0)
y_hc = kmeans.fit_predict(x)
```

```
[16]: y_hc
```

```
[16]: array([3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 0, 3, 3, 3, 3, 3, 3, 0, 0, 3,
          0, 3, 0, 3, 0, 0, 3, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
          0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 1, 0, 1, 1, 1, 1,
          1, 1, 1, 2, 1, 1, 1, 1, 1, 2, 2, 2])
```

```
[ ]:
```